Satyam Chandrakant Chatrola

+1 (973) 905-1864 sc10247@nyu.edu linkedin.com/in/satyamchatrola github.com/Nightshade14 Portfolio GCP Profile

EDUCATION

New York University — New York City, NY
Master of Science in Computer Science (CGPA: 3.61 out of 4.00) — *September 2023 – May 2025 (tentatively)*

Gujarat Technological University — Gujarat, India
Bachelor of Engineering in Computer Engineering (CGPA: 3.79 out of 4.00) — *June 2018 – June 2022*

SKILLS

Languages and DBs: **Python, SQL**, C, C++, Java, PostgreSQL, MySQL, MongoDB, Apache Solr, **Qdrant, Pinecone**.

AI and ML: SVMs, Gradient Boosted Trees, Clustering, **Computer Vision**, NLP, CNN, RNN, **Recommendation and Search Systems, Transformers**, Prompt Engineering, Generative AI (GANs, Diffusion models), **LLMs (RAG, PEFT, LoRA)**, Differential Privacy.

Data Science: NumPy, Pandas, Polars, Matplotlib, Seaborn, SciPy, Scikit-learn, **PyTorch, TensorFlow, Transformers, Optimum, Bitsandbytes, Ollama, Langchain**, Optuna, Tableau, Statistical Modeling, Stochastic Methods, **Hypothesis Testing, A/B Testing**.

MLE, MLOps and Data Eng.: **AWS SageMaker, ONNX, TensorRT**, MLflow, Evidently, Apache (Hadoop, Spark, Airflow, Kafka).

SWE and Others: **AWS, Google Cloud Platform (GCP)**, REST APIs, Git, **Docker, Kubernetes, CI/CD (CircleCI, GitHub Actions), Prometheus**, Flask, **Django, FastAPI, Redis, Pytest**, System Design, HTML, CSS, JavaScript, Node.js.

EXPERIENCE

Rapidops — Ahmedabad, India
Machine Learning Engineer + Data Scientist — *January 2022 – June 2023*

Face Authentication and Authorization System *(Python, PyTorch, YOLO, MTCNN, FaceNet, Triplet Loss, Qdrant, OpenCV)*
- Spearheaded a face authentication and authorization system with **4 environments and 2 access points** while **leading a team of 3**, from **data collection through on-premises deployment and monitoring**; enhanced productivity tracking and premises security.
- Architected a **scalable microservice** utilizing YOLO, MTCNN, fine-tuned FaceNet model, and Qdrant to achieve **low-latency** authentication and authorization on **live video streams**, delivering a **0.93 F1 score** without need to retrain for new individuals.

AI-Powered Search and Recommendation System *(Python, PyTorch, Hybrid Recommenders, Apache Solr, Docker, FastAPI)*
- Boosted **conversion rate** by **23%** and **click-through rate** by **42%** with **collaborative, content, and market-basket** recommenders.
- Engineered search & recommendations with **custom tagger**s & **LTR** techniques with **Apache Solr** to serve results in **25ms**.

Natural Language (English) to SQL query generation *(Python, TensorFlow, Keras, T5, BERT)*
- Researched and benchmarked **State of the Art (SOTA) AI models** generating SQL from Natural Language with **76% EMA**.
- Experimented with Transformers (**T5**, **BERT**) generating **73% Exact Match Accuracy (EMA)** with **36%** faster inference.

RESEARCH EXPERIENCE

Approaches to Type 2 Diabetes Mellitus Prediction with Machine Learning and Deep Learning
- Researched AI techniques for Type-2 Diabetes Mellitus classification with **95.8% precision, 95.8% recall**, and **99.4% specificity**.

PROJECTS AND OPEN-SOURCE CONTRIBUTIONS (with Embedded GitHub Repository URLs)

RAG Microservice: Research-mate chatbot *(Python, FastAPI, PyTorch, Transformers, AWS, Pinecone, Contextual RAG, Llamma 3.2 8B)*
- Engineered a **context-aware RAG**-based chatbot leveraging **Pinecone vector database**, enabling semantic search across **2,700 research papers** with **95% query relevance** and reducing response time to seconds.
- Optimized model performance with **Binary Quantization**, achieving **7x speedup** in **inference time** and **85% reduction** in **memory**.

Microservices webapp: LLM Essay Evaluator *(Python, PyTorch, ONNX, TensorRT, FastAPI, AWS, MLflow, Evidently, Frontend)*
- **Calibrated** Transformers (**BERT**) and LLMs (**GPT-2**) with **PEFT techniques** (quantization), cosine-annealed learning rate and warm-up to assess essays, attaining a **Kappa Score of 81.7%** and **surpassing the Benchmark** score by **5.7%**.
- Orchestrated **2 fault-tolerant microservices** & inferred with **ONNX models** and **TensorRT** via **asynchronous REST APIs**.

Open Source Project: mAIgic *(Python, OpenAI Function Calling, CircleCI, Pytest, MyPy, Ruff, uv)*
- Architected an AI-powered email management system leveraging **OpenAI's function calling API**, achieving **95% accuracy in task extraction** and automated Trello board updates, **reducing manual email processing time by 70%.**
- Engineered a **production-grade Python API** with **100% test coverage** through **CircleCI pipeline**, implementing comprehensive **static type checking** with mypy, and leveraged **SQLite**-based conversation tracking system.

Open Source Project: ETL pipeline migration to Spark *(Python, PySpark, AWS, MinIO, Databricks.Koalas)*
- Migrated **ETL pipeline** to **Spark** for **NYU's Open-Source** project with **160% speedup** in feature extraction for AI tasks.

CERTIFICATIONS

- Graduated from **Udacity's AWS Machine Learning Engineer Nanodegree** with top remarks.